Malaria Simulation (MaSim)
Version: 4.1.5.uganda

Robert Zupko

2023-08-30

Contents

1　Introduction

The malaria simulation, or MaSim, is a stochastic, individually-based model (or individually-based micro simulation) designed to investigate the evolution of antimalarial resistance by the *Plasmodium falciparum* in the presence of an antimalarial therapy given to a symptomatic individual, with transmission driven by a force of infection model. The simulation may be run as either a regionally-based model or a geographically-based model. While both approaches include the same levels of population heterogeneity, when run as a regionally-based model spatial heterogeneity (i.e., varying prevalence and transmission intensity) is limited whereas the geographically-based model allows for heterogeneity to be modeled at the configuration resolution of the simulation (e.g., 25 km^2).

1.1　Individuals

All individuals within the simulation represent humans that may become infected by the *P. falciparum* parasite, which leads to either a symptomatic or asymptotic infection. In the event of a symptomatic infection, individuals will seek treatment at the configured rate and will receive a therapy based upon the configuration provided to the simulation. Upon taking a therapy, a simplified pharmacokinetic and pharmacodynamic (PK/PD) is used to determine the blood concentration of the drug(s) in blood. The blood concentration, combined with the drug resistance profile for the *P. falciparum* clone(s) determines the killing rate of the parasite and reduction of blood parasitemia.

1.2　Infection by *P. falciparum*

While *P. falciparum* is transmitted by female *Anopheles* mosquitoes, as a modeling simplification a force of infection (FOI) model is used that presumes that transmission is driven in part by the total parasite load of all individuals within a given area. The likelihood that an individual is bitten is then determined by their individual biting attractiveness and a seasonal adjustment to account for the variable number of mosquitoes present throughout the year. All bites are presumed to be infectious and the probability of infection is determined by the individual's underlying immune response to *P. falciparum*.

1.3　*P. falciparum* Genotypes

Upon model initialization, the circulating genotypes of *P. falciparum* is determined by the configuration; however, when mutations are enabled, there is a small probability of a mutation occurring. While this mutation may confer an evolutionary advantage (e.g., resistance to an antimalarial), if accompanied by a fitness cost, the mutation may be lost or out competed by wild-type parasites without the mutation. Also included in the model is recombination due to interrupted feeding by a mosquito resulting in multiple clones being present in the gut. The probability for this occurring is based upon the configuration provided and the FOI of the circulating genotypes.

1.4　Policy Interventions

Within the simulation, drug policy interventions (e.g., a change in first-line therapies) are supported with limited support for non-therapeutic interventions (e.g., bed net distribution). Due to the lack of an underlying vector model, interventions targeting the *Anopheles* mosquitoes are not supported.

Simulation

2   Model Calibration

This chapter will introduce the reader to calibration of the simulation when spatial modeling is involved. The directions assume that work is being using a Linux workstation, or Windows workstation with Windows Subsystem for Linux (WSL) installed. Scripts referenced can typically be found in the PSU-CIDD-MaSim-Support repository and prior to running the scripts it will be necessary to add them to the `PATH`.

2.1   Geographic Data

By design the configuration of the simulation is flexible in terms what must be supplied as raster data and what can be supplied as a single value. In general, the following raster files are required to run the model spatially:

- *Beta*: The beta parameter determines, in part, how many additional infections occur as a result of a single infection and ultimately influences the *Pf*PR. These values must be calculated for each country based upon the inputs.
- Population: The model is initialized with the population values supplied in each of the cells in the raster. This file is mandatory when running spatially and can be generated by down sampling the spatial resolution a publicly available source while preserving the population count.
- *Pf*PR$_{2\text{-}10}$: While not directly used by the IBM, a raster of the *Pf*PR values is needed as part of the beta raster generation.

The following raster files a may or may not be needed depending upon the configuration and research goals:

- Access to treatment: In the event that access to treatment is variable depending upon the location within the country, it is necessary to supply a raster file defining the percentage of infected individuals that seek / receive treatment.
- Climate or malaria seasonality: If there are multiple seasonal variations to the malaria season in the country, it is necessary to supply a raster file that links the cell to the seasonal equation to be used.
- Political divisions: Since results are recorded in connection to a political district or region, a file is required to supply these values.
- Travel time or friction surface: Depending upon the movement algorithm selected, a travel time or friction surface is needed to ensure that individuals move in a reasonable fashion.

Once all of the GIS data has been acquired it and aligned, it is necessary to convert it to an Esri ASCII raster format (`asc` extension) to be loaded into the model.

2.1.1   Intitial Population

Typically the simulation is configured to use a constant crude birth rate during execution, which in turn requires that an *initial population raster* be produced. Assuming a reference population of 1,000 individuals, the initial population can be calculated by applying the equation:

$$p' = \frac{1000}{\left(1 + \frac{cbr}{1000}\right)^n}$$

Where *cbr* is the crude birth rate, n is the number of years prior to calibration year, and p' is the initial population. The cell multiplier can then be calculated as:

$$multiplier = \frac{1000 - p'}{1000}$$

The resulting *multiplier* can then be applied using the Raster Calculator functionality in geographic information system (GIS) software (e.g., ArcGIS Pro, GRASS GIS, or QGIS) where it is recommended that the floating point result be rounded up to the nearest integer value (i.e., Round Up function in ArcGIS Pro) to ensure that cells with a population always maintain a population. Once applied, the results can be validated by applying the equation:

$$projected = population * \left(1 + \frac{CBR}{1000}\right)^n$$

Where *projected* is the projected population n years after the initial *population* size for the given crude birth rate. Depending upon the number of digits used when applying the *multiplier* there may be some difference in the projected population versus the reference population and while ± 1 individual is accepted for large populations, the result can typically be corrected by increasing the number of digits in the *multiplier*. For cells with a small reference population (ex., < 10) in the focus should be on ensuring that a population remains in the cell since low population cells are subject to high variance in malaria dynamics during model execution.

2.2 Database Preparations

It is recommended that a new database be created for each country being studied. One of the limiting factors in the design of the database and IBM is that it presumes that the sequence of genotypes in each configuration file is the same. Once a new database has been created, the genotypes need to be loaded via the following command:

```
./MaSim -l -i [CONFIGURATION]
```

2.3 Beta Calibration

2.3.1 Setting Bounds for Beta Alignment

Once the GIS files have been prepared, work can proceed to generation of the beta parameters.

1. Begin by ensuring that all of the ASC format raster files are aligned correctly using `validateraster`
   1.1. If they are not, it will be necessary to determine why they are not aligned correctly.
   1.2. In general, rasters should be aligned to the population and PfPR$_{2\text{-}10}$ rasters, based upon which one has the most complete information.

2. Following alignment, the `generatebins` command should be run to generate the first calibration script to be run on the cluster this will set the bounds for when the betas are likely in relation to the PfPR$_{2\text{-}10}$ values for the country.
   2.1. First, run `generatebins` using the primary configuration file (e.g., `Studies/rwa-configuration.yml`).
   2.2. Login to the cluster and create a directory do the calibration from (e.g., `rwa-calibration`).
   2.3. Copy the following files to the cluster via `sftp` into the previously created directory:

   - From the country repository (e.g., PSU-CIDD-Rwanda):
     - `Studies/out/calibration.sh`
     - `Studies/Calibration/rwa-calibration.yml`
   - From the support repository:
     - `bash/calibrationLib.sh`

- `bash/manage.sh`
- `bash/runCalibration.job`
- `bash/template.job`
- `bash/population.asc`
- `bash/zone.asc`

2.4. Being the control script via the following command: `qsub runCalibration.job`
2.5. During calibration be sure to monitor for any errors, in general if replicate files are not deleted by the management script the error log should be checked.

3. Once all of the replicates have been run by `runCalibration.job` run the `createbetamap` command to load the calibration data from the database and create the first beta map.

2.3.2 Refining Beta Alignment

Once the bounds have been found it is necessary to reduce the epsilon values to an acceptable margin of error compared to the reference data. The is done by the iterating on the beta values using the `reduceepsilons` script until the sub-national annual rates approximate reference values.

1. Run the `createbetamap` command and note the maximum epsilon value output.
   1.1. If the value is outside of acceptable bounds, run the `reduceEpsilons` command using the maximum epsilon rounded down (e.g., if the maximum is 0.0314 then start with 0.01) and the same decimal place for the step value (i.e., 0.01 per the previous example).
   1.2. The script will generate `out/reduction.csv` and `out/script.sh` which needs to be uploaded to the cluster.
   1.3. Once the files have been uploaded run the job on the compute cluster.

2. Once all of the replicates have run on the cluster, run the `createbetamap` script.
   2.1. Note the maximum epsilon.
   2.2. Examine the map of epsilon values (`out/epsilons_beta.asc`) and note the distribution of high epsilon values.
   2.3. If they correspond to higher population areas repeat Step 1
   2.4. If they are concentrated largely in low population areas proceed to Step 3.

3. As the epsilon values are reduced, they need to be checked against the reference values by performing a validation run.
   3.1. Using the beta map generated and the status quo parameters for the country, run the model at scale.

3 Running the Simuation

3.1 Model Exeuction

When the simulation is executed, the operation can be broken into two distinct phases:

1. *Initialization* in which the simulation reads and verifies the YAML configuration file, and instantiates model data collector, population, and other relevant objects (see Figure 1).
2. *Execution* during which the simulation is executed by executing all scheduled operations and events for each time step that the simulation is configured to run (see Figure 2).

Figure 1: Typical model initialization sequence

Figure 2: Typical model time step sequence

3.1.1 Initalization

The main entry point for the simulation is via the `main` function, which first configures last chance error reporting (on Linux platforms), followed by parsing of the command line, configuration of logging, and model initialization (see Figure 1).

Model initialization proceeds by first reading the configuration file, which entails the sequential processing of YAML elements.[1] If the YAML file indicates that the simulation is based upon a raster model of geography, then the geographic data (e.g., population distribution, beta values, etc.) is loaded and used to prepare the internal location database (i.e., `location_db`). If a location model of geography is then present in the YAML file, an error will be produced. Likewise, if any errors are encountered during the processing of YAML notes, an error will be produced.[2]

Once the YAML has been successfully parse, the remaining model initialization proceeds in the following order:

1. The random seed for the simulation is generated and stored.
2. The model reporters are initialized sequentially in the order that they appear on the command line.
3. The event scheduler is initialized.
4. The initial treatment strategy is initialized.
5. The initial treatment coverage model is loaded.
6. The model data collector is initialized.
7. The initial population is initialized based upon the population described in the YAML file.
8. The movement model is initialized.
9. Initial malaria infections are introduced based upon the genotypes indicated in the YAML file.
10. Any environmental or population events that are configured in the YAML file are added.
11. Miscellaneous configuration occurs.

Typically the major operations involved in model initialization are preparing the initial population and infections; however, depending on the movement model and geography involved the one-time calculations involved in model movement may consume a significant amount of time.

3.1.2 Execution

Once initialization is complete, any pre-simulation operations are preformed as part of before run operations, followed by the primary simulation loop which is managed by the `Scheduler::run` function (see Figure 2).

Each time step of the simulation proceeds in the following manner:

1. If it is the first day of the month,
    1. The `monthly_report` function of each reporter supplied is executed in the order that they were added to the reporters list.
    2. The Model Data Collector resets any variables that are tracked by the month.
2. If it is the first day of the year,
    1. The Model Data Collector resets any variables that are tracked by the year.
3. Any scheduled population events are executed.
4. All daily population operations are executed,
    1. Infection events are preformed.

---

[1]While the order of elements in the YAML file does not matter, the simulation is highly sensitive to the order that configuration items are read and changes to the order of entries in the `src/Core/Config/Config.h` file should be carefully tested.

[2]If a superfluous node is encountered in the YAML file, no errors are produced; however, the lack of a required node will produce an error.

2. The birth event is preformed.
3. The circulation (i.e., population movement) is preformed.

5. Any scheduled individual events are executed.
6. End of time step operations are preformed.

As a result of the order of operations for the simulation, the model data is incremented before the monthly reporter is executed resulting in the date logged reflecting the start of the next month. However, because the Model Data Collector has not been reset, it will reflect all of the data of the previous month.

3.2 Troubleshooting

3.2.1 EOF encountered while reading data

This error can be caused by several sources, but if the file being read ends with a `.asc` extension, then it is most likely an Esri ASCII raster format file[3] (or ASC file) and the simulation encountered the end of the file before all of the expected data was read. Since the number of rows and number of columns is supplied at the start of an ASC file, the simulation attempts to read that number of cells and will produce this error if it runs out of data. The most likely cause of this error is edits to the ASC file that causes a mismatch between the number of cells and declared rows and columns.

3.2.2 MD5 hash collision

An MD5 hash collision error is produced by either the `DbReporter` or the `DbDistrictReporter` detects that YAML file has the same same file name and produces the same MD5 hash of an existing configuration in the database.[4] Under normal operations this error should be prevented by the unique constraint upon `sim.configuration.md5` and `sim.configuration.filename` although older instances of the database may contain the error. Correction of the error typically requires manual deletion of the duplicated hashes and all replicates associated with them via the `deleteReplicates.py` script in `PSU-CIDD-MaSim-Support/Python` followed by the addition of the constraint to the database:

```
ALTER TABLE IF EXISTS sim.configuration
    ADD CONSTRAINT configuration_md5_unique UNIQUE (md5, filename);
```

4 DxG Generator

The DxG Generator is used to calculate the efficacy of therapies in the simulation based upon the therapy's drug combination and the genotypes that are present in the YAML file for a given region. While this tool uses the same simulation code base, the executable is separate from the primary simulation (i.e., `masim`) and can be invoked as follows:

```
./DxGGenerator -i [input]
```

With the full scope of options supported:

```
-g                  Get the efficacies for a range of genotypes
-h / --help         Display the help menu
-i / --input        The YAML configuration to read
```

---

[3]https://desktop.arcgis.com/en/arcmap/latest/manage-data/raster-and-images/esri-ascii-raster-format.htm

[4]The `DbDistrictReporter` is an inheriting class of `DbReporter` so the error is logged as being produced by `DbReporter::prepare_configuration` regardless of which reporter is being used.

```
-t                      Get the efficacies for a range of therapies
--ec50                  EC50 for AS on C580 only
--iov                   AS inter-occasion-variability
--iiv                   AS inter-individual-variability
```

Note that the DxG Generator is very sensitive to the YAML configuration since the *full* configuration is still loaded prior to the efficacy calculations. As such it is recommended that a second, simplified (i.e., non-spatial) configuration be prepared for the DxG Generator. The tool is also only comparable with single therapies (i.e., `SFTStrategy`), which can be used in the configuration as follows:

```yaml
strategy_db:
  0:
    name: SFTStrategy
    type: SFT
    therapy_id: 0
initial_strategy_id: 0
```

Additionally, the DxG Generator was written with the assumption that `log_parasite_density_detectable` represents the lower detectable limit for a careful clinical study, typically 10 parasites per microliter of blood. However, if `log_parasite_density_detectable_pfpr` is set, then that value will be used instead. Since the *Pf*PR is driven by the more common detectable range of 50 to 100 parasites per microliter of blood, this will result in the efficacies returned being higher than expected. As such, it is recommended that the `log_parasite_density_detectable_pfpr` setting be deleted from the YAML configuration when it is prepared for the DxG Generator.

Development

5 Development

This chapter outlines the basic steps that are needed to build the simulation across various development environments. The guide is written from the standpoint of a "clean" workstation with no other dependencies in place.

5.1 Tool Chain Dependencies

5.1.1 Windows 10 / Windows Subsystem for Linux (WSL 1)

Step 1. Install a Linux sub-system of your choice (Ubuntu recommended, or Red Hat)

Step 2. Update the Linux sub-system

```
sudo apt update
sudo apt upgrade
```

Step 3. Install the build dependencies

```
sudo apt install build-essential
sudo apt install cmake
sudo apt install libgsl-dev
sudo apt install libyaml-cpp-dev
sudo apt install libfmt-dev
sudo apt upgrade
```

Step 4. Add PostgreSQL to the Apt Repository listing

Note that while this step is necessary when working from a new installation of Ubuntu, it may not be necessary, and you should check to see if the package can be found by first performing `apt search postgresql-10`

```
sudo sh -c 'echo "deb http://apt.postgresql.org/pub/repos/apt $(lsb_release \
  -cs)-pgdg main" > /etc/apt/sources.list.d/pgdg.list'
wget --quiet -O - https://www.postgresql.org/media/keys/ACCC4CF8.asc | sudo apt-key add -
sudo apt-get update
```

Step 5. Install the PostgreSQL and libpqxx dependencies

Note that this should be done from a directory that you are comfortable with git repositories being stored in.

```
sudo apt install postgresql-10
sudo apt install libpq-dev
sudo apt upgrade
git clone https://github.com/jtv/libpqxx.git
cd libpqxx
git checkout 7.0.0
./configure --disable-documentation
make
sudo make install
```

This may need to be done every time you connect to the VPN, so scripting the update may be in order.

5.1.2   Upgrading to WSL 2

WSL 2 offers some performance improvements which may be relevant during development so upgrading may be of interest. To do so the following needs to be performed, note that these steps will also enable WSL 2 if WSL 1 has not already been installed.

Open `PowerShell` with administrator permissions and enter the following:

```
dism.exe /online /enable-feature /featurename:Microsoft-Windows-Subsystem-Linux /all `
   /norestart
dism.exe /online /enable-feature /featurename:VirtualMachinePlatform /all /norestart
```

Once complete, restart the computer, after which the version needs to be set in `PowerShell`:

```
wsl --set-default-version 2
```

If you have already installed a distribution, updating is recommended. Otherwise, you can now install a WSL compliant distribution.

5.2   Building

5.2.1   Local WSL Builds

Before building the first time it is necessary to create the build directory within the local clone of the `PSU-CIDD-Malaria-Simulation` repository:

```
mkdir build
cd build
```

After which the following command can be run to build the simulation under WSL:

```
cmake -DCMAKE_BUILD_TYPE=DEBUG -DBUILD_WSL:BOOL=true ..
make -j 8
```

Generally it is recommended to create a `build.sh` script that runs the build commands.

5.2.2   Building on ICDS-ACI, Roar Collab

To build the simulation to run on the ICDS-ACI Roar Collab cluster the first time it is necessary to perform a number of configuration steps. After logging on to the interactive environment (`submit.hpc.psu.edu`) and cloning this repository, run `config.sh` which will prepare the build environment. As part of the process a build script will be created at `build/build.sh` that will ensure the environment is set correctly to compile the code base.

When the `config.sh` script is done running, note that it will recommend changes to the `.bashrc` which can be edited via `vi ~/.bashrc` and adding lines similar to the following:

```
# Configure run time environment
module use /storage/icds/RISE/sw8/modules
module load gsl
LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/storage/home/USERNAME/work/build_env/postgres/lib
```

Running the `config.sh` script will report the relevant lines at the end of the script.

5.3 Execution

5.3.1 Local Runs

The first step in performing a basic model check is load the genome data in the database, this must also be done whenever a new database is corrected:

```
cd build/bin
./MaSim -i ../../misc/input.yml -l
```

Once the genome data has been loaded the basic input file can be run as follows:

```
cd build/bin
./MaSim -i ../../misc/input.yml
```

Note that while care was taken in places to ensure the code is performant, the amount of RAM needed during execution can be quite high (ex., 32 GB or more). When the model is run in Linux environments where the necessary memory is not available, you may find that the program is killed without notice due to being out of memory.

5.3.2 Cluster Runs

The Roar Supercomputer Users' Guide providers a good overview for running single replicates on the cluster; however, when running batches it is recommended to script out the process. When replicates need to be run with a variety of settings (e.g., sensitivity analysis) some scripts present in PSU-CIDD-MaSim-Support under the `bash` directory can be used to parse a CSV formatted list of replicates to be run. In addition to the `calibrationLib.sh` file the support repository, the following files need to be created for this:

1) A runner script which will be queued on the cluster as a job, typically named `run.sh` or similar in project repositories:

```
#!/bin/bash
source ./calibrationLib.sh
runReplicates 'replicates.csv' '[USERNAME]'
```

2) The Portable Batch System (PBS) file that defines the job for `run.sh`:

```
#!/bin/bash

#PBS -A [ALLOCATION]
#PBS -l nodes=1:ppn=1:rhel7:stmem
#PBS -l pmem=4gb
#PBS -l walltime=120:00:00

#PBS -m ea

cd $PBS_O_WORKDIR
./run.sh
```

When defining the PBS file note the low memory usage (`pmem`) and high `walltime`. Since the job will only be responsible for running this script, only a limited amount of resources is needed. However, the total batch of jobs may run for quite some time, so the wall clock time is likely to be quite high.

3) The CSV file that defines the replicates to be run, where the first column is the PBS file for the replicate and the second column is the count:

```
bfa-slow-no-asaq.job,1
bfa-fast-no-asaq.job,1
bfa-rapid.job,1
```

While the `runReplicates` command executes, the number of jobs per user account is limited to the `LIMIT` defined in `calibrationLib.sh` (99 by default). When the limit is reached the script will sleep and periodically awaken to check to see if more jobs can be queued.

5.4 Development Tools

5.4.1 Isolating Segmentation Faults in Linux

When developing new functionality in the model it is sometimes necessary to isolate segmentation faults. One of the easier ways to do this is through the use of `gdb` in a Linux environment. First, compile the program with the `debug` flag set, this will ensure that there are debug symbols in the binary. Then use `gdb` to open the gdb console:

```
file bin/MaSim            # Load the specified executable
run -i ../input/sample.yml # Run the program with the following command line parameters

...                       # Program output omitted

bt                        # Generates the stack trace
```

5.4.2 Profiling

While there are many approaches to code profiling, as a quick way to get up and running, valgrind is recommended along with gprof2dot. Following installation, the simulation can be profiled using:

```
valgrind --tool=callgrind ./bin/MaSim - ../input/sample.yml
gprof2dot -f callgrind callgrind.out.* | dot -Tpng -o output.png
```

This will generate a PNG file containing a node graph of where most of the simulation's time is spent during execution along with the percentage of time spent in a given function.

Note that `valgrind` adds **considerable** overhead to the execution of the simulation, so it is highly recommended that profiling be limited to small simulations.

5.4.3 Valgrind under WSL 2 with CLion integration

When using CLion as an IDE, integration with Valgrind is possible. Presuming that WSL 2 has been enabled and `valgrind` has been installed, CLion simply needs to be informed of the path under WLS 2 via **File > Settings > Build, Execution, Deployment > Valgrind** and the path will typically be `/usr/bin/valgrind`.

5.5 Troubleshooting

5.5.1 WSL on the PSU VPN

One minor problem that may occur while on the PSU VPN is that `psu.edu` domains are not resolved correctly. This is usually due to the `/etc/resolv.conf` file not being updated correctly by WSL. To manually update the file, launch `PowerShell` and run:

```
Get-DnsClientServerAddress -AddressFamily IPv4 | Select-Object -ExpandPropert ServerAddresses
```

Then copy the addresses to `/etc/resolv.conf` so you have something similar to:

```
nameserver 192.168.1.1
nameserver 192.168.1.2
nameserver 192.168.1.3
search psu.edu
```

6 Database Infrastructure

Due to the complexity of data storage requirements for simulations involving multiple cells, a database is provided through the **dbreporter** reporter class. The schema is outlined below and was developed against PostgreSQL due to the high storage requirements. Note that the **notes** table presume the usage of the MaSimLIMS and can be removed if not needed with no impact upon the simulation.

6.1 Installation

The following guide is intended provide a walk through of how to get the simulation database up and running. While the document is kept as general as opposed, there may be some differences that you encounter due to your local IT requirements. The following guide assumes that the server is running a clean installation of Ubuntu 18.04 LTS 64-bit with ports 80, 443, and 5432 open.

6.1.1 Hardware Requirements

The specific requirements for the server are dependent in part upon the number of instances that will be connecting to it while a simulation is running. However, as a baseline the following is a reasonable starting point for a virtual machine:

- 4 CPUs
- 8 GB RAM
- 800 GB primary disk
- Ubuntu 18.04 LTS 64-bit

6.1.2 Installing PostgreSQL

1. Connect to the server

```
ssh [User]@[IP address] -p [port]
```

2. Set the default encoding on the server, the setting can be verified by running `locale`

districtmovement

PK	id
FK1,I1	replicateid timestep count source destination

configuration

PK	id
	yaml md5 name notes filename ncols nrows xllcorner yllcorner cellsize studyid
FK1,I1	

replicate

PK	id
FK1,I1	configurationid seed starttime endtime movement aggregationlevel

movement

PK	id
FK1,I1	replicateid timestep individualid source destination

location

PK,U1	id
FK1,I1,U1 U1	configurationid index x y beta district

notes

PK	id
	data user date studyid
FK1,I1	

study

PK	id
U1	name

monthlydata

PK	id
FK1,I1	replicateid dayselapsed modeltime seasonalfactor entrytime

therapyrecord

PK,FK1 PK PK	monthlydataid locationid therapyid
	success failure completed

monthlysitedata

PK,FK1 PK	monthlydataid location
	population clinicalepisodes treatments treatmentfailures eir pfprunder5 pfpr2to10 pfprall infectedindividuals nontreatment

monthlygenomedata

PK,FK2 PK,FK1 PK	monthlydataid genomeid location
	occurrences clinicaloccurrences occurrences0to5 occurrences2to10 weightedoccurrences

genotype
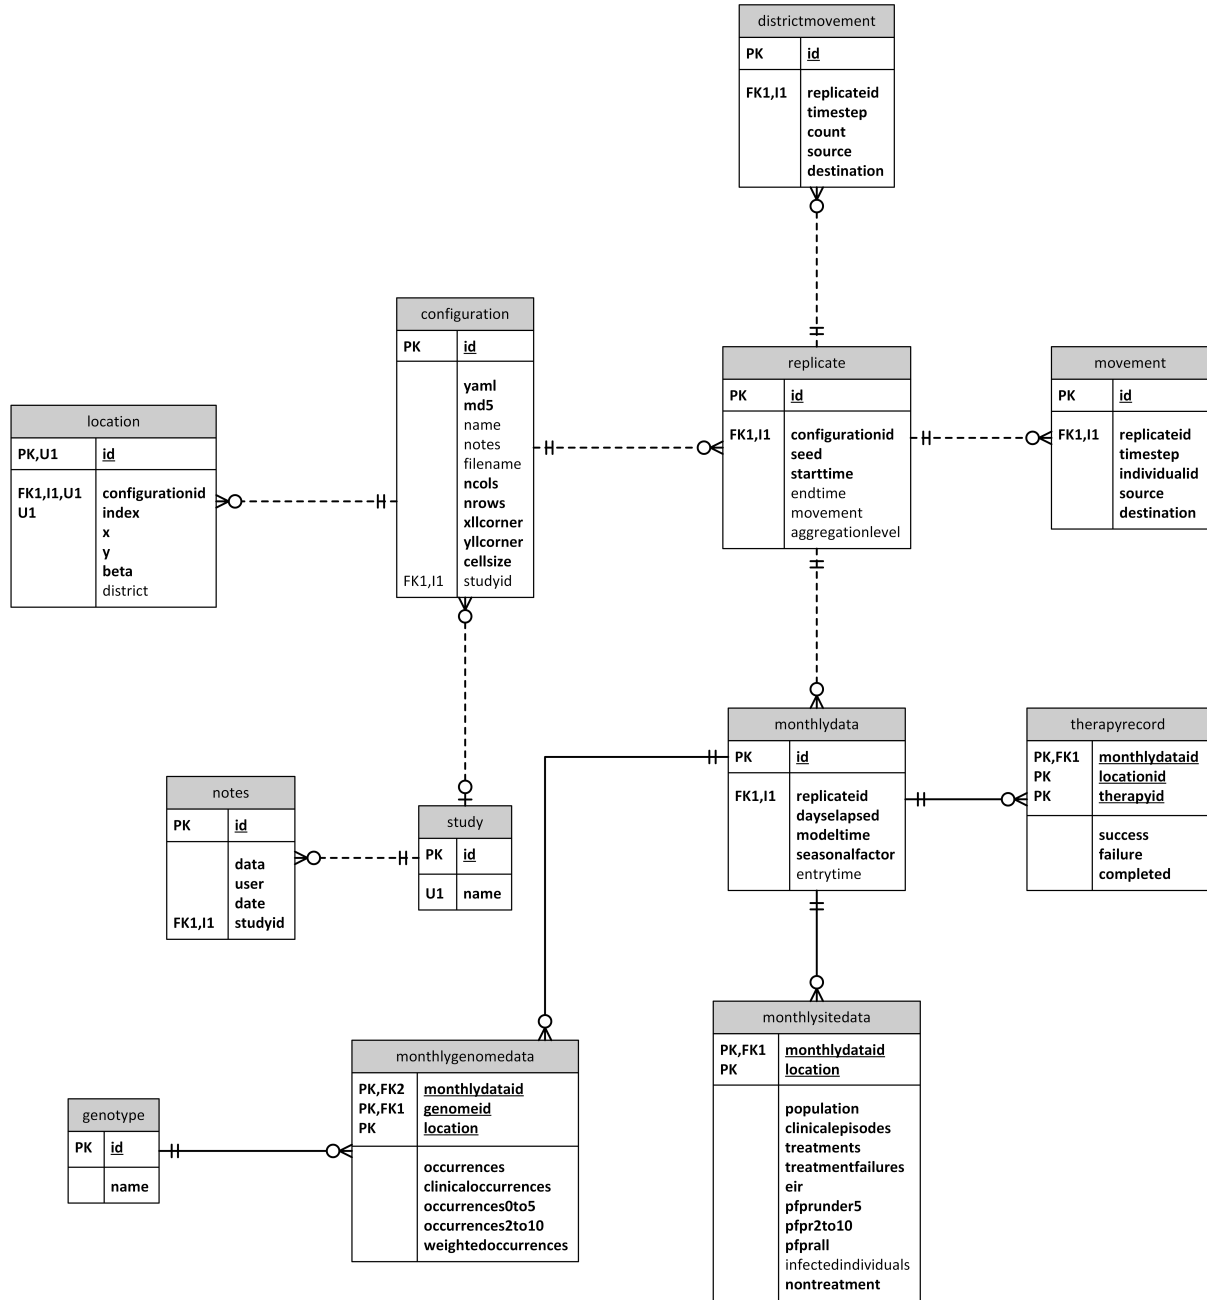
PK	id
	name

Figure 3: Database schema

```
export LANG=en_US.UTF-8
```

2. Install the GPG key and repository for PostgreSQL packages

```
sudo apt-get install wget ca-certificates
wget --quiet -O - https://www.postgresql.org/media/keys/ACCC4CF8.asc | sudo apt-key add -
sudo sh -c 'echo "deb http://apt.postgresql.org/pub/repos/apt/ \
  `lsb_release -cs`-pgdg main" >> /etc/apt/sources.list.d/pgdg.list'
```

3. Install PostgreSQL

```
sudo apt-get update
sudo apt-get install postgresql postgresql-contrib
```

4. Verify connection to PostgreSQL

```
sudo -u postgres psql postgres
postgres-# \conninfo
postgres-# \q
```

5. Create administrative user

Be sure to supply the password and select yes when asked if the user should be a super user.

```
sudo -u postgres createuser --interactive --pwprompt
```

6. Configure the server to listen for connections

This is done by first editing the file **/etc/postgresql/11/main/postgresql.conf** and updating the line **listen_addresses='*'**. Next, the file **/etc/postgresql/11/main/pg_hba.conf** needs to have the line **host all all 0.0.0.0/0 md5** added at the end. Note that this means that the server will listen to connections from *any* IP address. If this is not desired behavior, a more restrictive configuration should be used.

7. Enable service

```
sudo update-rc.d postgresql enable
```

8. Optional update for locale

Update default locale for the database template. This is necessary if you get an error that states "Encoding UTF8 does not match locale en_US. The chosen LC_CTYPE setting requires encoding LATIN1" from pgAdmin when creating a database

```
sudo -u postgres psql postgres

update pg_database set datistemplate=false where datname='template1';
drop database Template1;
create database template1 with owner=postgres encoding='UTF-8' lc_collate='en_US.utf8' \
  lc_ctype='en_US.utf8' template template0;
update pg_database set datistemplate=true where datname='template1';
\q
```

6.2 Installation of pgAdmin

The preferred way of installing pgAdmin on a clean Ubuntu installation is though the use of `apt`:

1. Configure the system for the pgAdmin APT repository.

```
sudo apt install curl
sudo curl https://www.pgadmin.org/static/packages_pgadmin_org.pub | sudo apt-key add
sudo sh -c 'echo "deb https://ftp.postgresql.org/pub/pgadmin/pgadmin4/apt/$(lsb_release \
  -cs) pgadmin4 main" > /etc/apt/sources.list.d/pgadmin4.list && apt update'
```

2. Install pgAdmin

```
sudo apt install pgadmin4-web
sudo /usr/pgadmin4/bin/setup-web.sh
```

6.2.1 Configuration of Apache for pgAdmin

Since newer version of pgAdmin are designed for web environments, configuration following installation may be limited editing the `pgadmin4.conf` configuration to reside at the server root:

```
WSGIDaemonProcess pgadmin processes=1 threads=25 python-home=/usr/pgadmin4/venv
WSGIScriptAlias / /usr/pgadmin4/web/pgAdmin4.wsgi

<Directory /usr/pgadmin4/web/>
    WSGIProcessGroup pgadmin
    WSGIApplicationGroup %{GLOBAL}
    Require all granted
</Directory>
```

The configuration can then be reloaded as follows:

```
sudo a2dissite 000-default.conf
sudo a2ensite pgadmin4.conf
sudo systemctl restart apache2
```

At this point you should be able to connect to the pgAdmin control panel at `http://[SERVER IP ADDRESS]`. Login to the control panel using the credentials supplied during configuration. One logged in, you should be able to add the localhost via "Add New Server" and proceed with administration of the databases using pgAdmin. Additional deployment information can be found on pgAdmin.org under Server Deployment.

7 Using the Database

7.1 Creation of Simulation Database

After logging into the pgAdmin control panel, start by creating the user `sim` and ensuring they have permissions to login to the database. This is the user that will be the simulation to write results to the database during model execution. Next, run the script `database.sql` which can be found under the `/database` directory of this repository.

7.2 Cloning databases

For the purposes of development or archiving it may be necessary to clone databases. The following SQL commands can be used from **psql** on the server to do so:

```sql
UPDATE pg_database SET datallowconn = false WHERE datname = 'masim';
CREATE DATABASE development WITH TEMPLATE masim OWNER sim;
UPDATE pg_database SET datallowconn = false WHERE datname = 'masim';
```

Demonstration

8 Demonstration

8.1 Introduction

The configuration included in the `demo` directory (`demo.yml`) is configured for a single cell and derives the majority of the settings from the Burkina Faso project, allowing it to emulate a realistic starting point for a model. In addition to the configuration file, there is also a basic raster file (`population.asc`) which informs the model on how to create a single cell with 1,000 individuals in it. The raster files used by the simulation are compliant with the Esri ASCII raster format and can be edited using a text editor, or appropriate geographic information system (GIS) software.

In order to use this demonstration code you first need to build the simulation using the guide outlined in Development, after which is recommended that you copy the binary (`masim`) and contents of the `demo` directory (i.e., `demo/demo.yml` and `demo/population.asc`) into the same location. The simulation can then be executed using the following command:

```
./MaSim -i demo.yml -r CellularReporter
```

Using the `CellularReporter` will generate a comma separated values (CSV) file in the directory containing model output for select parameters such as population, PfPR$_{2\text{-}10}$, and infected individuals to name a few. Running the simulation a second time will cause this file to be overwritten unless the job number (`-j`) switch is used. The `CellularReporter` is intended for models with only one location (i.e., cell) and will generate an error if multiple cells are present in the configuration.

8.2 Column Descriptions

The following tables are present in the output from the `CellularReporter`:

Table 1: Column format for the `CellularReporter`

Column	Description
DaysElapsed	Model days elapsed, should correspond to the length of months.
Population	Total population in the simulation.
PfPR2to10	The *Plasmodium falciparum* parasite prevalence for individuals aged 2 to 10 years.
TreatmentCoverage	The current treatment coverage in the simulation as an average for the under-5 and over-5 groups.
InfectedIndividuals	The number of individuals in the simulation who have *any* level of parasitemia.
ClinicalIndividuals	The number of individuals in the simulation who have a parasitemia level of clinical symptoms to manifest.
ClinicalU5	The number of newly clinical individuals under the age of 5.
ClinicalO5	The number of newly clinical individuals over the age of 5.
NewInfections	The number of new infections in the past month.
Treatment	The number of clinical individuals who received treatment.

Column	Description
NonTreatment	The number of clinical individuals who did not seek treatment.
TreatmentFailure	The number of clinical individuals who sought treatment and the treatment did not clear the infection in the configured amount of time.
ParasiteClones	The number of parasite clones present in the simulation, one individual may be infected by multiple clones.
Theta	The mean measure of immunity to the parasite in the simulation, zero means no immunity.
580yWeighted	The weighted number of occurrences of the 580Y genotype, such that the sum of all clones in an individual equals one.
508yUnweighted	The total number of occurrences of the 580Y genotype in the simulation.
Plasmepsin2xCopyWeighted	The weighted number of occurrences of the Plasmepsin double copy mutation, such that the sum of all clones in an individual equals one.
Plasmepsin2xCopyUnweighted	The total number of occurrences of the Plasmepsin double copy mutation in the simulation.

8.3   Example Console Output

```
[INFO] [default] MaSim version 4.0, experimental
[INFO] [default] Model initializing...
[INFO] [default] Read input file: demo.yml
[WARNING] [default] p_infection_from_an_infectious_bite used default value of 0
[INFO] [default] location_db appears to have been set by raster_db
[WARNING] [default] as_iov used default value of 0.2
[INFO] [default] Random initializing with seed: 1625174844985602
[INFO] [default] Starting day is 2007-01-01
[INFO] [default] Location count: 1
[INFO] [default] Population size: 1000
[INFO] [default] Model starting...
[INFO] [default] Perform before run events
[INFO] [default] Simulation is running
[INFO] [default] 17:27:24 - Day: 0
[INFO] [default] 2007-01-01 : turn mutation off
[INFO] [default] 17:27:25 - Day: 30
[INFO] [default] 17:27:25 - Day: 60
[INFO] [default] 17:27:25 - Day: 90

...

[INFO] [default] 17:28:19 - Day: 5400
[INFO] [default] 17:28:19 - Day: 5430
[INFO] [default] 17:28:19 - Day: 5460
[INFO] [default] Perform after run events
[INFO] [default] Model finished!
[INFO] [default] Final population: 1412
[INFO] [default] Elapsed time (s): 54.991
```

Appendicies

A   Genotype Information

Genotypes are encoded into the application via the input file which contains a YAML entry for `genotype_info`. This structure is organized as `loci` with the associated `alleles` that may be mutated during model execution. The following table outlines loci and alleles that are included in the sample configuration file.

Table 2: Genotype encoding scheme

Locus	Allele	Short Name	Description
PfCRT	K76	K	Chloroquine sensitivity, Amodiaquine sensitive, Lumefantrine resistance
	76T	T	Chloroquine resistance, Amodiaquine resistance, Lumefantrine sensitive
PfMDR1	N86 Y184 one copy of pfmdr1	NY–	Lumefantrine resistance, Amodiaquine sensitive
	86Y Y184 one copy of pfmdr1	YY–	Amodiaquine resistance
	N86 184F one copy of pfmdr1	NF–	lumefantrine resistance
	86Y 184F one copy of pfmdr1	YF–	Amodiaquine resistance, Lumefantrine resistance
	N86 Y184 2 copies of pfmdr1	NYNY	lumefantrine resistance, Amodiaquine sensitive, Mefloquine resistance
	86Y Y184 2 copies of pfmdr1	YYYY	Amodiaquine resistance, Mefloquine resistance
	N86 184F 2 copies of pfmdr1	NFNF	lumefantrine resistance, Mefloquine resistance
	86Y 184F 2 copies of pfmdr1	YFYF	Amodiaquine resistance, Lumefantrine resistance, Mefloquine resistance
K13 Propeller	C580	C	Artemisinin sensitive
	580Y	Y	Artemisinin resistance
Plasmepsin 2-3	Plasmepsin 2-3 one copy	1	Piperaquine sensitive
	Plasmepsin 2-3 2 copies	2	Piperaquine resistant
Hypothetical locus for multiple use	naïve	x	Experimental use in the model
	mutant	X	Experimental use in the model

The short name field requires additional note since genotype results generated by the model are based upon the short names. For example:

KNY–C1x

Indicates that the parasite has the K76 allele from the PfCRT locus (K), N86 Y184 one copy of pfmdr1 (NY–), C580 from K13 Propeller locus (C), Plasmepsin 2-3 one copy (1), and is a naïve copy of the hypothetical locus (x).

B   Technicalities

Due to the complex nature of the simulation, there are a number of technical decisions that may impact the model performance or results. This chapter contains the technical documentation concerning these decisions.

B.1   Random Numbers

A single random number generator is used for the entire life cycle of the simulation The simulation class `Random` acts as a wrapper for GNU Scientific Library Random Number Generation and Random Number Distributions. The simulation uses the Mersenne Twister 19937 generator (`gsl_rng_mt19937`), an implementation of the Mersenne twister (Matsumoto and Nishimura 1998), which produces 32-bit numbers with a state size of 19937 bits. Unless a seed is provided by the configuration, the random number generator instance is seeded by `std::random_device` which produces a non-deterministic, uniformly-distributed integer value.

B.2   Internal Events

Internal events are scheduled events with the simulation that are cannot be produced as a result of the YAML configuration. These are automatically created by the simulation when appropriate and may be of interest in understanding how the simulation works or how it can be modified.

B.2.1   `ProgressToClinicalEvent`

The `ProgressToClinicalEvent` is scheduled when an individual is infected by a parasite and the infection will progress to a clinical infection, and possible death. Upon the event triggering, the simulation first checks to see if the conditions for a clinical infection are still present (i.e., parasites still in their system) and that another clinical infection did not already trigger. If the conditions are acceptable then parasite density is increased to clinical levels, and the determination to seek treatment occurs. If the individual seeks treatment, then it is administered and record keeping takes place and the `TestTreatmentFailureEvent` is scheduled. Otherwise, recording keeping takes place and the individual immune response is replied upon to clear the infection. In the event that the dies due to the infection, the death is immediately recorded and any treatment is logged as a failure by `ReportTreatmentFailureDeathEvent` at the treatment failure testing day (`tf_testing_day`).

B.2.2   `ReportTreatmentFailureDeathEvent`

The `ReportTreatmentFailureDeathEvent` is produced by the `ProgressToClinicalEvent` when an individual receives treatment for malaria but dies. This event ensures that treatment failures are logged appropriately for the expected treatment time frame.

B.2.3   `TestTreatmentFailureEvent`

The `TestTreatmentFailureEvent` is created when an individual with clinical case of malaria elects to receive treatment and is scheduled to trigger on the configured treatment failure testing day (`tf_testing_day`), which is typically 28 days. Treatment failures are determined by testing to see if

the parasite that caused the clinical case has a parasite density above the detectable limit, as configured by the `log_parasite_density_detectable` field. The results of this test are logged for both the individual's location and the associated therapy.

C   Reporters

The simulation exposes a number of different reporters that can be specified via the command line switch `-r` and listed via the `--lr` switch. Reporters currently write data either to a text file (CSV or TSV formatted) or to the database back-end.

Primary reporters write extensive amounts of information from the simulation, in contrast, *secondary* reporters tend to be limited in either what is reported, or have limitations on their use.

C.1   Reporter Types

C.1.1   Database Reporter (`DbReporter`)

Primary Reporter
Output: Database

The `DbReporter`, short for Database Reporter, uses the PostgreSQL database as the back-end and respects the `record_genome_db` setting in the configuration file when determining if data genotype data should be stored. If genotype data is stored, it is aggregated at the cellular level and can result in large databases.

C.1.2   Database Reporter by District (`DbReporterDistrict`)

Primary Reporter
Output: Database

The `DbReporterDistrict` is a derivative of the `DbReporter`, except genotype data will be aggregated to the district level.

C.1.3   Genotype Carriers Reporter (`GenotypeCarriers`)

Specialist Reporter
Output: Database

The Genotype Carriers Reporter (`GenotypeCarriers`) is a specialist reporter that works in tandem with either of the database reporters and expects to be run after all other reporting has been completed, which can be done by invoking the reporter as the last option in the switch (e.g., `-r DbReporterDistrict,GenotypeCarriers`). When the reporter is triggered it will report count of the first allele mutation on the second locus (i.e., 580Y, 561H, or 469Y).

When the reporter initializes it will check for the presence of the `genotypecarriers` column in `sim.monthlysitedata` which is where the absolute count of genotype carriers will be stored during model executing. If the column is not present, it can be added using the following SQL:

```
ALTER TABLE sim.monthlysitedata
ADD genotypecarriers integer;
```

C.1.4 Therapy Record Reporter (`TherapyRecord`)

Specialist Reporter
Output: Database

The Therapy Record Reporter (`TherapyRecord`) is a specialist reporter that works in tandem with either of the database reporters and expects to be run at any point after the primary reporter has completed, which can be done by invoking the reporter following the main reporter in the switch (e.g., `-r DbReporterDistrict,TherapyRecord,GenotypeCarriers`). The reporter starts reporting data after the `start_collect_data_data` is reached and will report the total number of treatments that completed in the current month along with the total numbers of success or failures.[5] The data is recorded in the `sim.therapyrecord` table with the `locationid` field corresponding to either the cell id, or the district id, in accordance with the primary reporter used (i.e., `DbReporter` or `DbDistrictReporter`).

While the `sim.therapyrecord` will be created when the canonical template database is cloned, but may not exist in legacy databases. As such, if the reporter is passed as a parameter to the simulation, when the reporter initializes it will check for the presence of the `sim.therapyrecord` table and will call the simulation to terminate if the table is not present.

C.1.5 Monthly Reporter (`MonthlyReporter`)

Primary Reporter
Output: TSV File

The Monthly Reporter `MonthlyReporter` is the original reporter used by the simulation to report the simulation status and produces two tab separated values (TSV) file as described below.

C.1.6 Seasonal Immunity Reporter (`SeasonalImmunity`)

Specialist Reporter
Output: CSV File

The Seasonal Immunity Reporter (`SeasonalImmunity`) is a specialist reporter that scans all locations and individuals and aggregates summary data at the climatic zone level. This is useful for determining the role that seasonal transmission patterns has upon the genotype evolution and immune response to the parasite. The reporter is hard coded for the first allele mutation on the second locus, so while it reports 580Y in the output, 561H could be captured, depending upon model configuration.

C.2 Reporter Data Files

C.2.1 Monthly Reporter

The `MonthlyReporter` generates two data files in tab separated values (TSV) format with group separators indicated by the sentinel value `-1111` where *n* is a zero-indexed location id.

Table 3: **monthly_reporter_*n*.txt** - Summary data for the model generated at the end of each simulated month.

Block	Column Number(s)	Description
Summary	1	Model time, number of days elapsed
	2	Model time, calendar date as system time
	3 - 5	Model time, calendar date (Year, Month, Day)

---

[5] The total number of successes plus the total number of failures should always equal the total number completed.

Block	Column Number(s)	Description
	6	Seasonal factor
	7	Treatment coverage, probability to be treated (0 - 1)
	8	Treatment coverage, probability to be treated (0 - 10)
	9	Population size
	10	Group separator
EIR and PfRP	$1 + (n * 5)$	EIR by location per year
	$2 + (n * 5)$	Group separator
	$3 + (n * 5)$	Blood slide prevalence, PfPR < 5
	$4 + (n * 5)$	Blood slide prevalence, PfPR 2 - 2
	$5 + (n * 5)$	Blood slide prevalence, PfPr all
Infections by Location	1	Group separator
	$2 + n$	Number of new infections by location
Treatments by Location	1	Group separator
	$2 + n$	Number of treatments by location
Clinical Episodes by Location	1	Group separator
	$2 + n$	Number of clinical episodes by location
Genotype Frequency	1	Group separator
	. . .	See genotype frequency discussion

Table 4: **summary_reporter_*n*.txt** Summary that is generated after the model has completed execution.

Block	Column Number(s)	Description
Summary I	1	Random seed value
	2	Number of locations
	3	*Beta* value
	4	Population size
EIR and PfRP	$1 + (n * 5)$	EIR by location per year
	$2 + (n * 5)$	Group separator
	$3 + (n * 5)$	Blood slide prevalence, PfPR < 5
	$4 + (n * 5)$	Blood slide prevalence, PfPR 2 - 2
	$5 + (n * 5)$	Blood slide prevalence, PfPr all
Summary II	1	Treatment strategy id
	2	Percent treatment failures per year

C.2.2 Genotype Frequency

Genotype frequency is a complex entry that follows a similar output structure, but is generated differently depending upon the approach selected.

Genotype frequency, by weighted number of parasite-positive individuals

```
FOR EACH location
  FOR EACH state per person
    FOR EACH age class per person
      IF the person has no parasites THEN CONTINUE
      Update the total count
      Count the number of genotypes per person
      Adjust the count for the region
    END
```

27

```
    END
    // Formats the genotype per location as TSV
    FOR EACH genotype
      OUTPUT << weighted value << '\t'
    END
END
OUTPUT << "-1111\t"
// Format the summary genotype results as TSV
FOR EACH genotype
  OUTPUT << weighted value << 't'
END
OUTPUT << "-1111\t"
OUTPUT << Total count
```

D   MaSim Configuration File

The simulation uses YAML to load configuration settings for the simulation. While the simulation was previously forward compatible (i.e., older files would work with newer versions), with the transition to version 4.0 new nodes were added or deprecated that have resulted in a divergence between 3.x and 4.0 on wards.

As a matter of convention, the YAML key is generally indicated with **bold** text. The data type (e.g., integer, string, etc.) or the possible values are indicated in parentheses following the key name with any default indicated in **bold**. Generally this document is organized such that headings are organized by operational impact upon the simulation, followed by simple key-value pairs as the first entry in each section, followed by more complex entities as subheadings. Within subsections, the YAML keys should be in alphabetical order, although closely coupled keys (e.g., `number_of_age_classes` and `age_structure`) will break this pattern with the first key that should be read appearing first.

D.1   Model Operation

The following nodes govern how the model executes in terms of simulation execution.

connection_string (string) : The connection string for the PostgreSQL database that stores the simulation data.

days_between_notifications (integer) : The number of model days that should elapse between status updates to the console.

record_genome_db (Boolean) : Indicates that genome data should be recorded to the database when using the `DbReporter` reporter class. Note that recording genomic data to the database will cause the database to quickly inflate in size. It is recommended that this setting only be used when genomic data needs to be retrieved.

D.2   Model Configuration

The following nodes contain the settings for the simulation.

artificial_rescaling_of_population_size (double) : A scaling value that should be applied to the population size in a given location. Defaults to 1.0, but 0.25 is commonly applied when geospatial data is used that maps locations to current populations.

birth_rate (float) : The number of births per thousand individuals in the simulation, expressed as a decimal (i.e., 42 births per 1000 is entered as 0.042).

initial_age_structure (integer array) : Used to initialize the population structure at model initialization (time zero).

initial_seed_number (integer) : The seed value that should be used by the random number generator. The default value of zero (0) indicates that the seed will be generated at model execution time based upon the number of milliseconds since the Unix epoch.

number_of_age_classes (integer) : The size of the `age_structure` array.
age_structure (integer array) : An array of integer values that corresponds to the oldest age that defines a break in the age structure. This age structure is used for reporting and age-specific mortality calculations.
death_rate_by_age_class (float array) : A float array of values that corresponds to the all-causes death rate for the simulation withe same index correspondence as `age_structure`. Typically, supplied as a malaria adjusted value.
mortality_when_treatment_fail_by_age_class (float array) : A float array of values that corresponds to the death rate when treatment fails, using the same index correspondence as `age_structure`.

number_of_tracking_days (integer) : The number of days to take the total number of parasites in the population.

report_frequency (integer) : The number of model days that various reporters will save data, and data aggregation events will trigger.

start_collect_data_day (integer) : The number of model days that should elapse before data collection begins (e.g., number of clinical episodes, number of deaths, etc.)

start_of_comparison_period (date string, YYYY/mm/dd) : The calendar date upon which the simulation should start calculating the number of treatment failures (NTF), artemisinin monotherapy usage (AMU), and useful therapeutic life (UTL). Note that as of version 4.0 the AMU results are considered to be *deprecated* and will be removed at a later date.

starting_date (date string, YYYY/mm/dd) : The start date of the simulation, model days elapsed will be indexed from this date.
ending_date (date string, YYYY/mm/dd) : The end date of the simulation.

D.3   Simulation Geography

D.3.1   raster_db

This node contains data related to the spatial organization of the model to include population distributions and raster files.

Usage
The use of the `raster_db` node will override the use of the `location_db` and errors or inconsistencies will occur if both are used at the same time. All the raster files must have the same header information as defined in the Esri ASCII raster format and the same number of defined pixels. A `std::runtime_error` may be generated if the number of data pixels in a given raster exceeds, or is less than a previously loaded raster.

Note that while any arbitrary value may be used for the `NODATA_VALUE` it is recommended that the standard value of `-9999` be used.

```
raster_db:
  beta_raster: "beta.asc"
  district_raster: "district.asc"
  location_raster: "location.asc"
  population_raster: "population.asc"
  travel_raster: "travel.asc"
  ecoclimatic_raster: "ecoclimatic.asc"
  pr_treatment_under5: "pr_treatment_under5.asc"
```

```
pr_treatment_over5: "pr_treatment_over5.asc"

cell_size: 5
age_distribution_by_location: [[0.0]]
p_treatment_for_less_than_5_by_location: [0.0]
p_treatment_for_more_than_5_by_location: [0.0]

beta_by_location: [0.0]
```

beta_raster (string) : The beta parameter (float) used for each cell in the model, overrides the value set in `beta_by_location`.

district_raster (string) : Contains the canonical identification values (integer) that should be for each cell in the model. The simulation follows the same convention as ArcGIS Pro in that the identification values should be sequential from one to n and will verify this when loading the ASC file. Zero indexed districts are supported and consistency is checked (i.e., ten districts would be numbered 0 to 9).

location_raster (string) : May contain arbitrary numeric values and is used by the simulation to convert the X-Y coordinates to their linear values. Only needed if it is the *only* raster file provided.

population_raster (string) : Contains the number of individuals (integer) in each cell within the simulation.

travel_raster (string) : Contains the friction surface (float) that is used for determining the likelihood of travel to a given cell.

ecoclimatic_raster (string) : Contains the ecoclimatic zone value for each cell. The ecoclimatic zone value should be a zero-indexed integer that corresponds to the `seasonal_info` value to use.

pr_treatment_under5 (string) : Contains the access to treatment (float) for individuals under five years (60 months) for each cell in the simulation, overrides `p_treatment_for_less_than_5_by_location`.

pr_treatment_over5 (string) : Contains the access to treatment (float) for individuals over five years (60 months) for each cell in the simulation, overrides `p_treatment_for_more_than_5_by_location`.

cell_size (float) : The size of each cell along one axis, in kilometers.

age_distribution_by_location : An array of arrays which contains floating point values corresponding to the age distribution (as defined by `age_structure`).

p_treatment_for_less_than_5_by_location, **p_treatment_for_more_than_5_by_location** : (*Deprecated in Version 4.0*) Contains the access to treatment for individuals under or over (respectively) five years (60 months). If the array length is one, then the value supplied (float) is used for all locations in the model. Otherwise, each element in the array is presumed to correspond to a different location.

beta_by_location : (*Array deprecated in Version 4.0, single value recommended or raster*) Contains the beta value (float) to be used in the simulation. If the array length is one, then the value supplied (float) is used for all locations in the model. Otherwise, each element in the array is presumed to correspond to a different location.

D.3.2   location_db

To Be Written.

D.3.3   seasonal_info

This setting governs the malaria season in the model and operates differently depending upon the setting and version of the simulation.

D.3.3.1   Equation based

An equation based model of seasonal variation in transmission is provided where parameters must be fit to the following equation:

$$multiplier = base + \left(a \cdot sin^+ \left(\frac{b \cdot \pi \cdot (t - \phi)}{365}\right)\right)$$

Once the equation is fit, the following YAML is used for the configuration:

```yaml
seasonal_info:
  enable: true
  mode: equation
  equation:
    raster: false
    base: [0.0]
    a: [0.0]
    b: [0.0]
    phi: [0.0]
```

mode (**equation** | rainfall) : (*Optional*) Indicates the node that should be used for the seasonality, namely based upon the equation based model, or by using rainfall data. In the event that a value is not supplied, the simulation will default to the equation based model.

enable (true | false) : Enables or disables seasonality in the simulation.

raster (true | false) : Indicates that a raster file should be used to set the correct rate for each pixel. When true each index in the array corresponds to an array index used by the cells in the raster. Otherwise the first value is used for all pixels.

base, **a**, **b phi** (double array) : Arrays of one to n `double` values that indicate the the variables in the seasonal period equation.

D.3.3.2   Rainfall Based

The rainfall based approach assumes that a model has already been fit and the adjustment values are stored in a CSV file.

```yaml
seasonal_info:
  enable: true
  mode: rainfall
  rainfall:
    filename: filename.csv
    period: 365
```

mode (equation | **rainfall**) : Required in order to load the rainfall data.

filename (string) : The CSV file that contains the adjustment that should be applied to the beta. Each adjustment should be supplied on a single line in the file.

period (integer) : The period of time before the pattern in the CSV file should repeat, generally 365 days is expected.

D.3.4   spatial_model

The `spatial_model` setting defines movement model that transfers individuals between model cells. All spatial models used in the simulation have the same basic definition structure:

```yaml
spatial_model:
  name: "ModelName"
  ModelName:
    paramter: value
```

where the `name` may be of the type `Marshall`, `Wesolowski`, or `WesolowskiSurface` which each have their own configuration values.

D.3.4.1 Marshall Movement Model

The Marshall movement model is based upon the gravity model described in Marshall et al. (2018) which presumes that the probability of a trip is defined by the proportional probability of movement from i to j such that $P(j|i) \propto N_j^\tau k(d_{i,j})$ where the kernel is defined by $k(d_{i,j}) = \left(1 + \frac{d_{i,j}}{\rho}\right)^{-\alpha}$ from the following configuration:

```
spaital_model:
  name: "Marshall"
  Marshall:
    tau: 1.0
    alpha: 1.0
    log_rho: 1.1
```

tau (double) : The calibrated value for τ.
alpha (double) : The calibrated value for α.
log__rho (double) : The calibrated $log_{10}(\rho)$ value.

D.3.4.2 Wesolowski Movement Model

THe Wesolowski movement model is based upon the gravity model described in Wesolowski et al. (2015) where the amount of travel to N_{ij} is defined by $N_{ij} = \frac{pop_i^\alpha pop_j^\beta}{d(i,j)^\gamma}\kappa$ from the following configuration:

```
spaital_model:
  name: "Wesolowski"
  Wesolowski:
    kappa: 1.0
    alpha: 1.0
    beta: 1.0
    gamma: 1.0
```

kappa (double) : The calibrated value for κ.
alpha (double) : The calibrated value for α.
beta (double) : The calibrated value for β.
gamma (double) : The calibrated value for γ.

D.3.4.3 Wesolowski Movement Model with Travel Surface

This movement model is similar to the generic Wesolowski movement model, except a travel surface penalty is applied to N_{ij} such that $N'_{ij} = \frac{N_{ij}}{(1+t_i+t_j)}$ using the `travel_raster` loaded as part of the `raster_db` and the following configuration:

```
spaital_model:
  name: "WesolowskiSurface"
  WesolowskiSurface:
    kappa: 1.0
    alpha: 1.0
    beta: 1.0
    gamma: 1.0
```

kappa (double) : The calibrated value for κ.
alpha (double) : The calibrated value for α.
beta (double) : The calibrated value for β.
gamma (double) : The calibrated value for γ.

D.4   Individual Immunity and Infection Response

allow_new_coinfection_to_cause_symtoms (*true* | false) : Flag to indicate if an asymptomatic host that is bitten and infected by a new parasite clone may present with new symptoms. Note the spelling of `symtoms` in the configuration.

transmission_parameter (float) : Governs how likely malaria is to be transmitted to a naive individual in the sporozoite challenge.

D.4.1   parasite_density_level

The `parasite_density_level` setting contains several sub-values that govern individual behavior or state due to the total number of parasites that are present in the individual's blood stream. When setting the parasite density for the detectable levels note that 10 per μl is the middle of the bounds for detection under Giemsa-stained thick blood film under laboratory conditions (4 - 20 parasites/μl) while 50 per μl is the lower bounds for detection under field conditions (50 - 100 parasites/μl) (Wongsrichanalai et al. 2007). Generally, a higher detection limit for the PfPR will require a higher transmission for a given PfPR than a lower detection level.

```
parasite_density_level:
  # corresponds to 100 total parasites (0.00002 per ul)
  log_parasite_density_cured:        -4.699

  # corresponds to 50,000 total parasites (0.01 per ul)
  log_parasite_density_from_liver:   -2.000

  # corresponds to 1,000 parasites per microliter of blood
  log_parasite_density_asymptomatic:   3

  # corresponds to 20,000 parasites per microliter of blood (total 10^11)
  log_parasite_density_clinical:      4.301

  # corresponds to 2,000 parasites per microliter of blood (total 10^10)
  log_parasite_density_clinical_from:   3.301

  # corresponds to 200,000 parasites per microliter of blood (total 10^12)
  log_parasite_density_clinical_to:     5.301

  # corresponds to 10 parasites per microliter of blood
  log_parasite_density_detectable:      1.000

  # corresponds to 50 parasites per microliter of blood
  log_parasite_density_detectable_pfpr: 1.699

  # corresponds to 2,500 parasites per microliter of blood
  log_parasite_density_pyrogenic:       3.398
```

log_parasite_density_cured (double) : When an individual is considered to be **cured** of a specific parasite colony.

log_parasite_density_from_liver (double) : Governs the lower bound for the number of parasites following the initial infection.

log_parasite_density_asymptomatic (double) : Threshold at which an individual is asymptomatic of malaria.

log_parasite_density_clinical (double) : Governs the upper bound for the number of parasites following the initial infection.

log_parasite_density_clinical_from (double) : Governs the lower bound of parasites that an individual may be inflicted with when progressing to clinical via the `ProgressToClinicalEvent` event.

log_parasite_density_clinical_to (double) : Governs the upper bound of parasites that an individual may be inflicted with when progressing to clinical via the `ProgressToClinicalEvent` event.

log_parasite_density_detectable (double) : Sets the threshold for the number of parasites that an individual may have present in their blood when tested to check if the prescribed treatment failed.

log_parasite_density_detectable_pfpr (double) : Sets the threshold for the number of parasites that an individual may have present in their blood when tested to see if a *detectable* level is presented. This value is used to inform calculations for the PfPR in the simulation.

log_parasite_density_pyrogenic (double) : (**UNUSED**) Sets the threshold for when fever may present as a symptom.

D.4.2   immune_system_information

The `immune_system_information` node contains parameters that are used to control the response of the individual immune system and is one of the mechanisms by which the simulation can be calibrated to match a given country.

```
immune_system_information:
  # Immune function parameters
  b1: 0.00125
  b2: 0.0025

  # Duration of infection parameters
  duration_for_naive: 300
  duration_for_fully_immune: 60

  # Population initialization parameters
  mean_initial_condition: 0.1
  sd_initial_condition: 0.1

  # Probability bounds for clinical symptoms
  max_clinical_probability: 0.99

  # Immunity acquisition parameters
  immune_inflation_rate: 0.01
  age_mature_immunity: 10
  factor_effect_age_mature_immunity: 0.3

  # Immunity function parameters
  immune_effect_on_progression_to_clinical: 12
  midpoint: 0.4
```

b1 (double) : Rate of immune function increase when parasitaemic
b2 (double) : Rate of immune function decrease when not parasitaemic
duration_for_naive (double) : Duration, in days, of infection when naive.
duration_for_fully_immune (double) : Duration, in days, of infection when fully immune.

mean__initial__condition (double) : Mean initial immune function of population at initialization.
sd__initial__condition (double) : Standard deviation of initial immune function of population at initiation.
max__clinical__probability (double) : Maximum probability of clinical symptoms as a result of a new infection.
immune__inflation__rate (double) : Yearly age-dependent faster acquisition of immunity between ages 1 to 10.
age__mature__immunity (double) : Age at which the immune function is mature, i.e., age at which the immune acquisition model switches from child to adult.
factor__effect__age__mature__immunity (double) : Adjustment to the curve of immune acquisition under the age indicated by `age_mature_immunity`, parameter kappa in supplement to Nguyen et al. (2015).
immune__effect__on__progression__to__clinical (double) : Slope of the sigmoidal probability versus immunity function, parameter z in supplement to Nguyen et al. (2015).
midpoint (double) : Adjusts the midpoint of the slope of the sigmoidal probability versus immunity function, parameter z in supplement to Nguyen et al. (2015).

D.5   Treatments

The following nodes govern how drugs and treatments are managed by the simulation.

tf__testing__day (integer) : The number of days following the administration of a therapy that a individual should be tested for treatment failure.

D.5.1   drug__db

This setting is used to configure the various drugs used in the configuration and is structured as an array of drugs, which contain the specific setting for that particular compound. **Note** that while the simulation assumes that the compounds will be ordered from 0 to *n*, it is the responsibility of the user to ensure that they are assigned and ordered correctly.

```
drug_db:
  0:
    name: "ART"        # Artemisinin
    half_life: 0.0
    maximum_parasite_killing_rate: 0.999
    n: 25
    age_specific_drug_concentration_sd: [0.4,0.4,0.4,0.4,0.4,0.4,0.4,0.4,0.4,
                                         0.4,0.4,0.4,0.4,0.4,0.4]
    age_specific_drug_absorption: [0.7,0.7,0.85,0.85,0.85,0.85,0.85,0.85,0.85,
                                   0.85,1.0,1.0,1.0,1.0,1.0]
    mutation_probability: 0.005
    affecting_loci: [2]
    selecting_alleles: [[1]]
    k: 4
    EC50:
      ..0..: 0.75
      ..1..: 1.2
```

name (string) : The name of the compound.
half__life (double) : The compound's half-life in the body, in days.
maximum__parasite__killing__rate (double) : The percentage of parasites that the compound will kill in one day if an individual has the highest possible drug concentration.
n (integer) : The slope of the linear portion of the concentration-effect curve.
age__specific__drug__concentration__sd (double array) : The actual drug concentration, per individual,

that will be drawn from a normal distribution with a mean of one and this standard deviation.

age_specific_drug_absorption (double array) : The percentage of the drug that is absorbed into the bloodstream, based upon the age of the individual. When not supplied the default value is one for all age groups.

mutation_probability (double) : The probability that exposure to the drug will result in a mutation in the parasite to resist it.

affecting_loci (integer array) : The index of the loci of alleles where drug resistance may form (see genotype_info).

selecting_alleles (integer matrix) : The index of the alleles where drug resistance may form (see genotype_info).

k (double) : Controls the change in the mutation probability when drug levels are intermediate. For example, k=0.5 is a simple linear model where mutation probability decreases linearly with drug concentration; whereas k=2 or k=4 are a piecewise-linear model where mutation probability increases from high concentrations to intermediate concentrations, and then decreases linearly from intermediate concentrations to zero.

EC50 (array of key-value pairs) : The drug concentration which produces 50% of the parasite killing achieved at maximum-concentration, format is a string that describes the relevant genotypes (see genotype_info), followed by the concentration where 1.0 is the expected starting concentration.

D.5.2   therapy_db

This setting is used to define the various therapies that will be used in the simulation and two variations are supported: simple therapies that consist of one or more drugs (defined using the the `id` from the `drug_db`) given over a number of days, and complex therapies that consist of one or more therapies (defined using the `id` of the previously defined therapy) given over a regimen.

```
therapy_db:
  # Artemisinin combination therapy (ACT) - artemether-lumefantrine (AL), three days
  0:
    drug_id: [0, 1]
    dosing_days: [3]
  # ACT - AL, one day
  1:
    drug_id: [0, 1]
    dosing_days: [1]

  # Complex therapy, AL dosed three days, one day off, with one final dose (3-1-1)
  2:
    therapy_ids: [0, 1]
    regimen: [1, 5]

  # Artemisinin combination therapy (ACT) - artemether-lumefantrine (AL),
  # five days with specified compliance
  3:
    drug_id: [0, 1]
    dosing_days: [5]
    # Probability that an individual will complete exactly this many days of treatment
    pr_completed_days: [0.1, 0.1, 0.2, 0.2, 0.4]
```

D.5.2.1   Simple Therapies

drug_id (integer array) : One or more integers that correspond to the defined identification numbers (i.e., array index) in the `drug_db`.

dosing_days (integer) : The number of days that the drug combination should be given for.

pr_completed_days (float array) : (*Optional*) The probability that the individual will comply with the course of treatment where 1 indicates they will always take it; otherwise, $0 < n < 1$ is the probability that they will stop on that day. In the event that the field is not supplied then it is assumed that the individual will always comply with the therapy.

D.5.2.2 Complex Therapies

Complex therapies consist of multiple simple therapies that are dosed over several days and may contain gaps in the dosing. Prior to Version 4.1.1 patient compliance with the therapy was determined by the `p_compliance` which generally assumed full compliance with the course of treatment. With complex therapies, noncompliance is currently not support and such configurations will produce an error.

therapy_ids (integer array) : One or more integers that correspond to the defined therapies.
regimen (integer array) : A one-index list of the days that the corresponding therapy should be given.

D.5.3 strategy_db and initial_strategy_id

This setting describes the treatment strategy applied when an individual presents with a clinical case of malaria and typically consists of at least two strategies: an initial 'baseline' that is consistent with the national first-line therapy and one or more interventions that represent the various drug policy approaches that are being applied. Following model initialization and burn-in following the strategy defined in `initial_strategy_id`, a new strategy can be employed using the `change_treatment_strategy` which takes effect upon execution of the event.

Within the `strategy_db` the indexing should begin at zero although the `initial_strategy_id` allows for any defined strategy to be used. The `type` field used for defining the strategies comes from the `IStrategy::StrategyTypeMap` and are thus case-sensitive. Failure to supply a valid strategy will result in the `StrategyBuilder` logging a `WARNING` before returning `nullptr`.

```
strategy_db:
  0:
    name: Baseline
    type: MFT
    therapy_ids: [ 0 ]
    distribution: [ 1 ]
  1:
    name: ExampleMFT
    type: MFT
    therapy_ids: [ 1, 2 ]
    distribution: [ 0.5, 0.5 ]
  2:
    name: ExampleCycling
    type: Cycling
    therapy_ids: [ 1, 2 ]
    cycling_time: 90
initial_strategy_id: 0
```

initial_strategy_id (integer) : The integer index of the treatment strategy to use from day zero of the simulation.

D.5.3.1 Cycling Strategy

A cycling strategy calls for the rotation between two or more single first-line therapies (e.g., artemether–lumefantrine and artesunate–amodiaquine) on a strict fixed schedule (e.g., every 90 days as opposed to

quarterly). As implemented the simulation supports an unlimited number of days and will return to the first therapy defined in `therapy_ids`. Bookkeeping related to when cycling events occur are automatically managed by the simulation as part of the `Model::daily_update` function.

name (string) : A string name used to refer to the treatment in console output and logs.
type (string) : Always `Cycling`
therapy_ids (integer array) : The therapies to deploy, based upon the index of the therapies defined in `therapy_db`.
cycling_time (integer) : The number of days that a therapy is used before switching.

D.5.3.2 Multiple First Line Therapies (MFT)

The multiple first-line therapies (MFT) approach calls for two or more therapies to be deployed nationally with an individual randomly receiving one at time of treatment, which they then use for the entire course of treatment. As a matter of convention, when a single first-line therapy is deployed, the MFT strategy *should* still be used in the configuration over the deprecated `SFT` (single first-line therapy) strategy and the `MFTStrategy` class works correctly when only a single therapy is defined.[6]

name (string) : A string name used to refer to the treatment in console output and logs.
type (string) : Always `MFT`
therapy_ids (integer array) : The therapies that may be given, based upon the index of the therapies defined in `therapy_db`.
distribution (float array) : The probability that an individual will receive the therapy corresponding to the same index where the total of all values provided *must* equal one.

D.6 Policy Interventions

While most of the policy interventions can be implemented using the therapies deployed, and switching them via events such as `change_treatment_strategy`, some of the more complex strategies require more in-depth configuration or are closely coupled with how the simulation operates.

D.6.1 Regular administration of prophylactic therapy (version 4.1.3, untested)

The regular administration of prophylactic therapy, or RAPT protocol, is a speculative approach to malaria control that calls for individual to take an artemisinin combination therapy (ACT) periodically without the presentation of clinical symptoms for malaria. The protocol presumes an individual will remember the month to take then next ACT, and at some point during that month a check will be performed to see if they take the therapy. The probability is based upon the probability that an individual in their age group will seek treatment and the probability of compliance with the RAPT protocol (i.e., $Pr(RAPT) = Pr(Treatment) \cdot Pr(Compliance)$). In the event that the individual already took an ACT in the past 28 days (determined by checking for `TestTreatmentFailureEvent`), they will not take the ACT regardless.

Implementation of this protocol requires that events be scheduled for at model initialization, although the point at which the individuals start taking the ACTs is determined by the configuration. Due to the computational overhead involved with the RAPT protocol, the simulation execution time is longer. If the `rapt_config` entry is not present in the configuration file, the event will not be enabled within the simulation.

NOTE that the RAPT protocol is highly experimental and the event has not been extensively tested.

```
rapt_config:
  period: 12
  therapy_id: 1
  compliance: 0.7
```

---

[6]The `SFT` strategy is preserved for use with the DxG Generator, but should not be used in simulations.

```
    age_start: 18
    start_day: 2022/08/18
```

period (int): the interval, in months, between RAPT doses.
therapy_id (int): corresponds to id of the treatment defined in the `therapy_db` to take.
compliance (float): the probability that the individual will comply with the RAPT protocol.
age_start (int): the age, in years, that the individual will start checking for compliance with the RAPT protocol.
start_day (date string, YYYY/mm/dd): the date at which the RAPT protocol will take effect, after this date, individuals will start checking for compliance.

D.7   Genotype Information

D.7.1   genotype_info

To Be Written.

D.8   events

This setting is used to list the various events that will be loaded and run during the model. The `name` field dictates which event will be parsed and all the data for the `info` field following will be provided to the loader function.

D.8.1   annual_beta_update_event

The annual beta update event increases or decreases the beta for each cell in the model using the formula $beta' = beta + (beta \cdot rate)$ and clamps the lower bounds for the beta a zero.

```
events:
  - name: annual_beta_update_event
    info:
      - day: 2020/10/26
        rate: -0.025
```

day (date string, YYYY/mm/dd) : the date when the update will first occur.
rate (float) : the rate of change for the beta for the cells.

D.8.2   annual_coverage_update_event

The annual coverage update event increases the coverage for each cell in the model by reducing the coverage gap by a fixed value provided by rate. If the model is run for a long enough period of time, the presumption should be that the coverage will reach 100%.

```
events:
  - name: annual_coverage_update_event
    info:
      - day: 2020/09/02
        rate: 0.025
```

day (date string, YYYY/mm/dd) : the date when the update will first occur.
rate (float) : the rate of reduction in the coverage for the cells.

D.8.3 change_treatment_strategy

Change the treatment strategy that is currently deployed in the simulation to the one defined in `strategy_db` with the given index. Simple error checking is performed upon model initialization and the `strategy_id` must be within bounds of the `strategy_db`.

```
events:
  - name: change_treatment_strategy
    info:
      - day: 2024/1/1
        strategy_id: 1
```

day (date string, YYYY/mm/dd) : The date when the update will first occur.
strategy_id (integer) : The index of the strategy defined in `strategy_db`.

D.8.4 importation_periodically_random_event

Over the course of the month indicated (January:1 - December:12) introduce the infection into the model at a random location, selected by a draw that is weighted by the population. Each day an infection is introduced based upon a uniform draw against `count / [days in month]`. Once started, this event continues until model termination.

```
events:
  - name: importation_periodically_random_event
    info:
      - day: 2021/08/01
        genotype_id: 1
        count: 10
        log_parasite_density: 4.301
      - day: 2021/09/01
        genotype_id: 2
        count: 20
        log_parasite_density: 3.0
```

day (date string, YYYY/mm/dd) : The first day of the month for which the event will occur.
genotype_id (int) : The id of the genotype to be introduced.
count (int) : The number of cases to be introduced in the month.
log_parasite_density (double) : the log density of the parasite to be imported.

D.8.5 introduce_mutant_event

On the specified date, find infected individuals and force the parasite genotype from the given wild type to mutation specified (e.g., C580 to 580Y). This operation will fill the difference between the input fraction and the current frequency of the genotype in the population. **Note** that while this is a one time event, it is recommended that the event be invoked multiple times prior to any policy interventions acting upon a given mutation frequency.

```
events:
  - name: introduce_mutant_event
    info:
      - day: 2021/12/27
        district: 9
```

```
    fraction: 0.01
    locus: 2
    mutant_allele: 1
```

day (date string, YYYY/mm/dd) : The model date when this event should occur.
district (int) : The district id for where the mutation event should occur.
fraction (int) : The target frequency of the mutation.
locus (int) : The genotype database locus index of the desired mutation.
mutant_allele (int) : The genotype database allele index of the mutation allele that will be applied.

D.8.6 introduce_mutant_raster_event

Similar to the `introduce_mutant_event` in that infected individuals will be switched from teh wild type to the mutant genotype indicated; however, the selection of individuals is based upon the raster provided as opposed to a district number.

This event follows the "fail fast" paradigm and as part of event initialization during model initialization the raster is loaded and checked to ensure that the cells indicated correspond to populated locations. Determination of cell locations uses the same algorithm as all other raster processing in the simulation, and as a result simulated cells must be labeled as either zero (no mutations) or one (mutations).

```
events:
  - name: introduce_mutant_raster_event
    info:
      - day: 2021/12/27
        raster: locations.asc
        fraction: 0.01
        locus: 2
        mutant_allele: 1
```

day (date string, YYYY/mm/dd) : The model date when this event should occur.
raster (string) : An ASC file indicating the locations where the mutations should take place. A cell value of zero indicates no mutations, while a value of one indicates that mutations should occur.
fraction (int) : The target frequency of the mutation.
locus (int) : The genotype database locus index of the desired mutation.
mutant_allele (int) : The genotype database allele index of the mutation allele that will be applied.

D.8.7 turn_off_mutation

Turn off all mutations in the model, recommended during the model burn-in.

```
events:
  - name: turn_off_mutation
    info:
      - day: 2020/09/04
```

day (date string, YYYY/mm/dd) : the date when the event will occur.

D.8.8 turn_on_mutation

Turn on all mutations in the model, or the mutations for individual drugs.

```
events:
  - name: turn_on_mutation
    info:
      - day: 2020/09/04
        drug_id: 0
        mutation_probability: 0.005
```

day (date string, YYYY/mm/dd) : the date when the event will occur.
drug__id (integer) : the id of the drug, as defined in the `drug_db` or `-1` to apply the value to all drugs.
mutation__probability (float) : the mutation probability to use.

D.8.9   update__beta__raster__event

Update all beta values in the simulation to those in the raster file.

```
events:
  - name: update_beta_raster_event
    info:
      - day: 2023/02/24
        beta_raster: "beta_2023.asc"
```

day (date string, YYYY/mm/dd) : the date when the event will occur.
beta__raster (string) : The beta parameter (float) used for each cell in the model, overrides the current value in the cell.

D.8.10   update__ecozone__event

Update all the cells matching the original ecozone to the new ecozone.

```
events:
  - name: update_ecozone_event
    info:
      - day: 2021/01/23
        from: 0
        to: 1
```

day (date string, YYYY/mm/dd) : the date when the event will occur.
from (integer) : the id of the ecozone, defined in `seasonal_info`, that is the original ecozone.
to (integer) : the id of the new ecozone, defined in `seasonal_info`, that will be applied to matching cells.

References

Marshall, John M., Sean L. Wu, Hector M. Sanchez C., Samson S. Kiware, Micky Ndhlovu, André Lin Ouédraogo, Mahamoudou B. Touré, Hugh J. Sturrock, Azra C. Ghani, and Neil M. Ferguson. 2018. "Mathematical Models of Human Mobility of Relevance to Malaria Transmission in Africa." *Scientific Reports* 8 (1): 7713. https://doi.org/10.1038/s41598-018-26023-1.

Matsumoto, Makoto, and Takuji Nishimura. 1998. "Mersenne Twister: A 623-Dimensionally Equidistributed Uniform Pseudo-Random Number Generator." *ACM Trans. Model. Comput. Simul.* 8 (1): 3–30. https://doi.org/10.1145/272991.272995.

Nguyen, Tran Dang, Piero Olliaro, Arjen M Dondorp, J Kevin Baird, Ha Minh Lam, Jeremy Farrar, Guy E Thwaites, Nicholas J White, and Maciej F Boni. 2015. "Optimum Population-Level Use of Artemisinin Combination Therapies: A Modelling Study." *The Lancet Global Health* 3 (12): e758–66. https://doi.org/10.1016/S2214-109X(15)00162-X.

Wesolowski, Amy, Wendy Prudhomme O'Meara, Nathan Eagle, Andrew J. Tatem, and Caroline O. Buckee. 2015. "Evaluating Spatial Interaction Models for Regional Mobility in Sub-Saharan Africa." *PLOS Computational Biology* 11 (7): e1004267. https://doi.org/10.1371/journal.pcbi.1004267.

Wongsrichanalai, Chansuda, Mazie J. Barcus, Sinuon Muth, Awalludin Sutamihardja, and Walther H. Wernsdorfer. 2007. "A Review of Malaria Diagnostic Tools: Microscopy and Rapid Diagnostic Test (RDT)." *The American Journal of Tropical Medicine and Hygiene* 77 (6_Suppl): 119–27. https://doi.org/10.4269/ajtmh.2007.77.119.